Multiple AI Review Board

SQUOOSH | Praneel Pothukanuri | November 2, 2026

Build a startup idea review demo

Praneel, build a usable full-stack prototype where someone enters a startup idea and gets feedback from AI agents acting as specialists relevant to that business. The application must generate a different specialist panel for different ideas, rather than reuse generic agents with new names. A multiple AI review board means these specialists assess the idea independently, challenge one another's reasoning and combine their findings into practical feedback.

Generate the right specialist panel

Capture the idea, target customer, problem, proposed solution and any supporting evidence. Generate a panel of three to five specialist agents, each with a clear role, relevant expertise, review mandate and explanation of why the idea needs that perspective. Let the user inspect and adjust the panel before running it. For example, an ecommerce idea might need merchandising and fulfillment specialists; a developer-tool idea might need developer experience and infrastructure specialists. The specialist mix and questions must change meaningfully with the idea.

Connect agents to different AI providers

Keep specialist roles separate from the underlying model. Let the user assign a provider and model to each agent through a shared connector interface. Implement adapters for two different AI providers, such as OpenAI, Anthropic or Gemini, plus a deterministic mock adapter. Use the same structured response contract across providers. The complete demo must run without API keys; clearly distinguish simulated responses from live calls. Live API use is optional, and no paid access or company data is required.

Review the idea and combine the feedback

Each specialist first reviews the idea without seeing the others' answers. Require feedback tied to its expertise: strengths, risks, missing assumptions, questions and a proposed validation experiment. Cite supplied evidence IDs where available and label assumptions or unknowns. Run one bounded challenge and response round, then use a separate coordinator to summarize priorities, disagreements and next experiments. Explain when evidence is insufficient. Do not invent market facts or treat consensus as proof of demand; this is decision support, not an investment recommendation.

Make the flow usable and reliable

Show idea input, editable panel, per-agent progress, specialist feedback, discussion and a combined review. Save the idea, panel configuration, outputs and status so a user can reopen or rerun a review. Show provider/model and latency; label any token usage or cost estimates. Add timeouts, at most one retry and a configurable call/token cap. Show failures explicitly and explain their effect on the combined review. Validate response schemas and evidence IDs, treat input text as data rather than instructions, and keep keys server-side and out of logs.

Demo walkthrough

Show two contrasting startup ideas producing meaningfully different panels and specialist-specific feedback. Include a thin-evidence case, a provider failure and reopening a saved review. Demonstrate tests for panel structure, provider-response normalization, malformed output, invalid citations, timeouts and budget stops. A live or short recorded prototype demo is enough. Use code and a brief README to explain architecture, two relevant technical or research sources, trade-offs, limitations and AI coding assistance. Choose your stack. Authentication, billing and browser automation are outside scope.